

# CIS 606 Machine Learning, Spring 2013

---

Lecturers: Wei Lee Woon and Zeyar Aung

## Lecture 13

# Ensemble Learning

Original Source:

[www.cs.cornell.edu/Courses/cs4700/2008fa/PPT/CS4700-EL.ppt](http://www.cs.cornell.edu/Courses/cs4700/2008fa/PPT/CS4700-EL.ppt)

(by Prof. Carla P. Gomes, Computer Science Dept., Cornell University)

# Ensemble Learning

---

So far – learning methods that learn a **single hypothesis**, chosen from a hypothesis space that is used to make predictions.

**Ensemble learning** ▪ select a **collection (ensemble) of hypotheses** and **combine their predictions**.

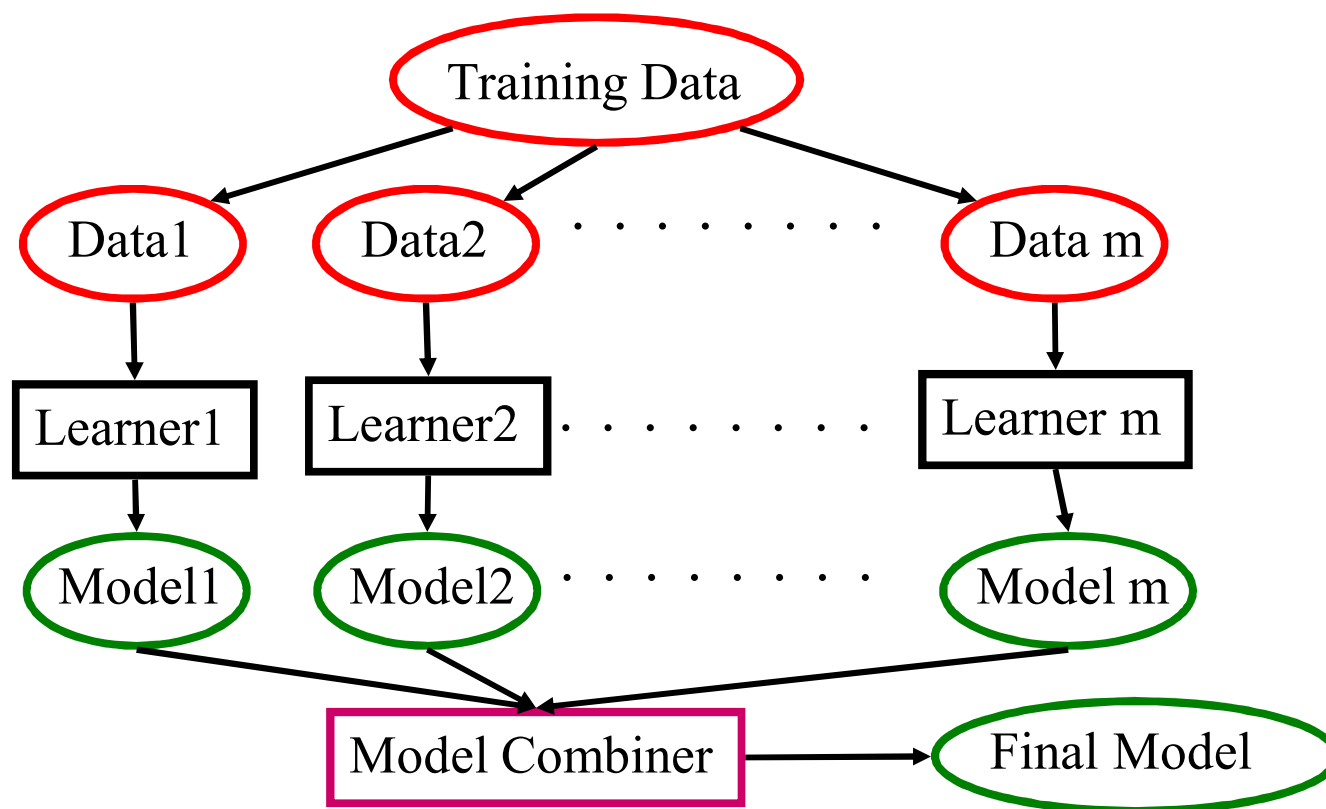
Example 1 - generate 100 different decision trees from the same or different training set and have them **vote on the best classification** for a new example.

**Key motivation:** reduce the **error rate**. Hope is that it will become much more **unlikely that the ensemble of** will misclassify an example.

# Learning Ensembles

Learn multiple alternative definitions of a concept **using different training data** or **different learning algorithms**.

**Combine decisions** of multiple definitions, e.g. using **weighted voting**.



# Value of Ensembles

---

## “No Free Lunch” Theorem


















































- No single algorithm wins all the time!

When combining multiple **independent** and **diverse decisions** each of which is **at least more accurate than random guessing**, random errors cancel each other out, **correct decisions are reinforced**.

Examples: Human ensembles are demonstrably better

- How many jelly beans in the jar?: Individual estimates vs. group average.
- Who Wants to be a Millionaire: Audience vote.

# Example: Weather Forecast

Reality							
1							
2							
3							
4							
5							
Combine							

# Intuitions

---

## Majority vote

Suppose we have 5 completely independent classifiers...

- If accuracy is 70% for each
  - $(.7^5) + 5(.7^4)(.3) + 10(.7^3)(.3^2)$
  - **83.7% majority vote accuracy**
- 101 such classifiers
  - **99.9% majority vote accuracy**

**Note: Binomial Distribution:** The probability of observing  $x$  heads in a sample of  $n$  independent coin tosses, where in each toss the probability of heads is  $p$ , is


$$P(X = x|p, n) = \frac{n!}{x!(n-x)!} p^x (1 - p)^{n-x}$$

# Ensemble Learning

---

Another way of thinking about ensemble learning:

- way of **enlarging the hypothesis space**, i.e., the ensemble itself is a hypothesis and the **new hypothesis space is the set of all possible ensembles constructible from hypotheses of the original space.**



**Increasing power of ensemble learning:**

Three linear threshold hypothesis  
(positive examples on the non-shaded side);  
Ensemble classifies as positive any example classified  
positively by all three. **The resulting triangular region** hypothesis  
is not expressible in the original hypothesis space.

# Different Learners

---

Different learning **algorithms**

Algorithms with different choice for **parameters**

Data set with different **features**

Data set = different **subsets**



# Homogenous Ensembles

---

Use a single, arbitrary learning algorithm but **manipulate training data** to make it learn multiple models.

- $\text{Data1} \neq \text{Data2} \neq \dots \neq \text{Data } m$
- $\text{Learner1} = \text{Learner2} = \dots = \text{Learner } m$

Different methods for changing training data:

- Bagging: Resample training data
- Boosting: Reweight training data

In **WEKA**, these are called **meta-learners**, they take a learning algorithm as an argument (**base learner**) and create a new learning algorithm.

---

# Bagging

# Bagging

---

Create **ensembles** by “*bootstrap aggregation*”, i.e., repeatedly **randomly resampling the training data** (Brieman, 1996).

**Bootstrap**: draw  $N$  items from  $X$  with replacement

## Bagging

- Train  $M$  learners on  $M$  bootstrap samples
- Combine outputs by voting (e.g., **majority vote**)

Decreases error by **decreasing the variance** in the results due to ***unstable learners***, algorithms (like decision trees and neural networks) whose output can change dramatically when the training data is slightly changed.

# Bagging - Aggregate Bootstrapping

---

Given a standard training set  $D$  of size  $n$

For  $i = 1 \dots M$

- Draw a sample of size  $n^* < n$  from  $D$  uniformly and with replacement
- Learn classifier  $C_i$

Final classifier is a vote of  $C_1 \dots C_M$

Increases classifier stability/reduces variance

---

# Boosting

# Strong and Weak Learners

---

**Strong Learner** ▪ Objective of machine learning

- Take labeled data for training
- Produce a classifier which can be *arbitrarily accurate*

**Weak Learner**

- Take labeled data for training
- Produce a classifier which is *more accurate than random guessing*

# Boosting

---

**Weak Learner:** only needs to generate a hypothesis with a training accuracy greater than 0.5, i.e.,  $< 50\%$  error over any distribution

## Learners

- Strong learners are very difficult to construct
- Constructing weaker Learners is relatively easy

Questions: Can a set of **weak learners** create a single **strong learner** ?

**YES ù**

**Boost weak classifiers to a strong learner**

# Boosting

---

Originally developed by computational learning theorists to guarantee performance improvements on fitting training data for a *weak learner* that only needs to generate a hypothesis with a training accuracy greater than 0.5 (Schapire, 1990).

Revised to be a practical algorithm, AdaBoost, for building ensembles that empirically improves generalization performance (Freund & Shapire, 1996).

## Key Insights

Instead of sampling (as in bagging) re-weight examples!

Examples are *given weights*. At each iteration, a new hypothesis is learned (*weak learner*) and the *examples are reweighted* to focus the system on examples that the most recently learned classifier got wrong.

Final classification based on *weighted vote of weak classifiers*




# Adaptive Boosting

---

Each rectangle corresponds to an example,  
with **weight proportional to its height**.

Crosses correspond to **misclassified** examples.

Size of decision tree indicates **the weight of that hypothesis** in the final ensemble.



# Construct Weak Classifiers

---

## Using Different Data Distribution

- Start with **uniform weighting**
- During each step of learning
  - **Increase weights** of the examples which are **not correctly learned** by the weak learner
  - **Decrease weights** of the examples which are **correctly learned** by the weak learner

## Idea

- Focus on difficult examples which are not correctly classified in the previous steps

# Combine Weak Classifiers

---

## Weighted Voting

- Construct **strong classifier** by **weighted voting of the weak classifiers**

## Idea

- Better weak classifier gets a larger weight
- Iteratively add weak classifiers
  - Increase accuracy of the combined classifier through minimization of a cost function

# Adaptive Boosting: High Level Description

---

$C = 0$ ; /\* counter\*/

$M = m$ ; /\* number of hypotheses to generate\*/

1 Set same weight for all the examples (typically each example has weight = 1);

2 While ( $C < M$ )

2.1 Increase counter  $C$  by 1.

2.2 Generate hypothesis  $h_C$ .

2.3 Increase the weight of the misclassified examples in hypothesis  $h_C$

3 Weighted majority combination of all  $M$  hypotheses (weights according to how well it performed on the training set).

Many variants depending on how to set the weights and how to combine the hypotheses. ADABOOST ▪ quite popular!!!!

# Performance of Adaboost

---

Learner = Hypothesis = Classifier

Weak Learner:  $< 50\%$  error over any distribution


M number of hypothesis in the ensemble.

If the **input learning is a Weak Learner**, then **ADABOOST** will return a **hypothesis that classifies the training data perfectly for a large enough M**, boosting the accuracy of the original learning algorithm on the training data.

**Strong Classifier:** thresholded linear combination of weak learner outputs.

# Restaurant Data

---



Decision stump: decision trees with just one test at the root.

# Restaurant Data

